Language-Gated GR-ConvNet: Steering a Compact Grasp CNN with Frozen CLIP Patch–Text Similarity

Minh-Ngoc Luong
Hanoi University of Science and Technology
ngoclm004@gmail.com

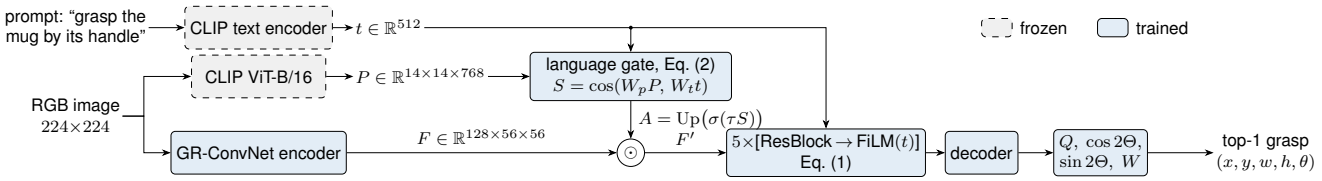


Figure 1. **Language-Gated GR-ConvNet.** Blue: trained (2.9M parameters); CLIP [5]: frozen. The prompt enters through a spatial gate A and through FiLM [4] after every residual block. C0 replaces $W_t t$ by a learned vector q.

Abstract

Language-driven grasp detection places a grasp on the object part named by a free-form prompt. The proposed method steers the GR-ConvNet grasp CNN with FiLM and a language gate*: a spatial map of cosine similarities between frozen CLIP patch tokens and the prompt embedding. On the annotated pairs of the official LGD test split of Grasp-Anything++, the model reaches 68.6% seen and 57.7% unseen success (two-seed mean), against 54.8% and 40.7% with FiLM alone. A control gate that receives the same CLIP features but not the prompt scores 11.6 and 11.8 unseen points lower in the two seeds ($p < 10^{-7}$), so prompt-dependent gating gives 9.4 of the 16.0-point gain in harmonic mean over FiLM. Code:* https://github.com/Ngoc-LM/language-gated-grasp.

1. Introduction

A parallel-jaw grasp [1] is a rotated rectangle $g = (x, y, w, h, \theta)$. Language-driven grasp detection conditions g on a prompt such as "grasp the knife by its handle", so the model must also locate the named part. On Grasp-Anything++ (1M images, 10M prompts), the reference method LGD [8] is a diffusion model. This work tests whether a single-pass grasp CNN can instead be steered by language. The contributions are (i) a *language gate* that re-weights GR-ConvNet [2] features by CLIP [5] patch–text similarity (Fig. 1); and (ii) a controlled ablation (text-free gate, two seeds, paired tests) that separates CLIP visual features from prompt-dependent gating, with open code.

2. Related Work

Generative grasp CNNs [2, 3] regress dense grasp maps. CLIPort [6] and CROG [7] use CLIP for manipulation and grasping, and LGD [8] denoises grasps under text guidance on data built with foundation models [9]. Pixel–text score maps from CLIP-style encoders drive dense prediction [10]; the language gate uses one as a spatial prior in a grasp CNN.

3. Method

Grasp maps. As in [2], the network predicts four 224×224 maps: quality Q, $\cos 2\Theta$, $\sin 2\Theta$ and opening W. Targets fill the central third of each ground-truth rectangle. The grasp is read at the most interior pixel of the maximum plateau of the smoothed Q, with $\theta = \frac{1}{2}\,\mathrm{atan2}(\sin 2\Theta, \cos 2\Theta)$ and $h = w/2$.

FiLM. The encoder maps the image to $F \in \mathbb{R}^{128\times56\times56}$, followed by five residual blocks and a decoder. With t the normalised CLIP text embedding, each residual block k is followed by

$$F \leftarrow \big(1 + \Delta\gamma_k(t)\big) \odot F + \beta_k(t), \tag{1}$$

where $\Delta\gamma_k, \beta_k \in \mathbb{R}^{128}$ come from a zero-initialised linear layer, so each FiLM layer is the identity at initialisation.

Language gate (proposed). FiLM applies the same channel modulation at every pixel, so it can favour one object among several only indirectly. The language gate therefore compares the 14×14 post-LayerNorm patch tokens P of a frozen CLIP ViT-B/16 with the prompt and re-weights the encoder output:

$$\begin{aligned} S_{ij} &= \cos(W_p\, p_{ij},\, W_t\, t), \\ A = \mathrm{Up}_{56}\big(\sigma(\tau S)\big), &\qquad F' = F \odot (1 + A), \end{aligned} \tag{2}$$

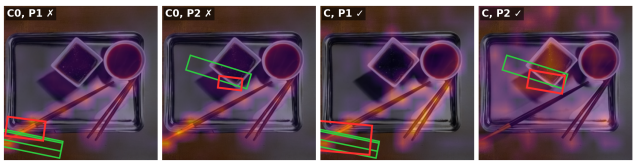


Figure 2. **Gate maps** (heat), top-1 grasps (red), GT (green); unseen image, seed 0. P1: "Hold chopsticks at its pair of sticks"; P2: "Take hold of bottle on its container". The C0 gate ignores the prompt; the C gate follows it. Selected as the largest C-gate change where only C is correct for both prompts.

	Model	Seen	Unseen	H	ms
A	GR-ConvNet (RGB only)	27.3	27.8	27.5	4.7
B	+ FiLM(t)	$54.8_{\pm 2.7}$	$40.7_{\pm 0.6}$	$46.7_{\pm 1.4}$	14.1
C0	B + text-free gate	$63.3_{\pm 0.5}$	$46.0_{\pm 1.9}$	$53.3_{\pm 1.1}$	19.6
C	B + language gate	$\mathbf{68.6}_{\pm 0.4}$	$\mathbf{57.7}_{\pm 1.8}$	$\mathbf{62.7}_{\pm 1.3}$	19.7
D	C + gate supervision	70.8	57.7	63.6	21.5

Table 1. **Success rate (%) on the annotated test split.** B, C0, C: two-seed mean±half-range. ms: batch-1 T4 latency incl. CLIP.

where $p_{ij} \in \mathbb{R}^{768}$ is a patch token, W_p and W_t are learned affine maps to $\mathbb{R}^{256}$, τ is a learned scale (initially 10), and Up_{56} bilinearly upsamples $S \in \mathbb{R}^{14\times14}$ to 56×56. The residual form $1+A$ only amplifies prompt-matching regions. The **C0 control** replaces $W_t t$ with a learned vector q: its gate sees the same CLIP features but not the prompt (FiLM still receives t), and it has 131k fewer parameters than C. **Loss:** Smooth-ℓ_1 summed over the four maps. Variant D adds $0.5\,\mathrm{BCE}(\tau S, M)$, M the pooled ground-truth Q mask.

4. Experiments

Protocol. Training uses 35,927 Grasp-Anything++ scenes (183,559 prompt–grasp pairs) and validation 734 scenes. Testing follows the official LGD seen/unseen split, restricted to annotated pairs: 1,376 seen and 601 unseen from 727 of 803 object keys. Test scenes, and training pairs whose object name is an unseen test name, are excluded from training; a word search finds unseen names in 5,332 training prompts, all as parts of other objects (e.g. a spoon's "bowl"), and 40 of the 41 seen names occur in training. A top-1 grasp is correct if its IoU exceeds 0.25 and its angle error is below 30° for some ground-truth grasp of the prompted part; H is the harmonic mean of seen and unseen success. All runs use AdamW (lr 10^{-3}, weight decay 10^{-4}, 500 warm-up steps, cosine), batch 32 and 6 epochs on one T4; checkpoints are selected on 1,000 validation pairs.

Results. FiLM (B) raises H by 19.2 points, CLIP features behind the text-free gate (C0) add 6.6, and letting the gate read the prompt (C) adds 9.4. Seed for seed, C exceeds C0 by 11.6/11.8 unseen and 6.3/4.4 seen points; paired McNemar tests give $p < 10^{-3}$ in all four comparisons, and scene-resampled 95% bootstrap intervals exclude zero. The unseen gain of C0 over B varies between seeds ($+7.8$, $p < 10^{-3}$; $+2.8$, $p = 0.21$), whereas the C0→C gain does not. C is correct for two prompts on different objects in 54 of the 118 multi-object test scenes, C0 in 40 (seed 0, Fig. 2). Prompt-dependent gating thus gives 69% of the B→C gain on unseen objects, 59% in H and 38% on seen ones. D does not differ significantly from C (seed 0; $p \geq 0.06$), so C is kept.

Cost and limitations. C adds 0.33M trained parameters to B and a frozen 149M-parameter CLIP (13.0 ms per frame with a cached prompt) and trains on 3.6% of the images. A exceeds the 18 unseen reported for GR-ConvNet with CLIP in [8], so LGD's 48/42 are not compared.

5. Conclusion

The gate raises the H of a FiLM-conditioned GR-ConvNet from 46.7 to 62.7; prompt-dependent gating gives 9.4 points.

References

[1] Yun Jiang, Stephen Moseson, and Ashutosh Saxena. Efficient grasping from RGBD images: Learning using a new rectangle representation. In *2011 IEEE International conference on robotics and automation*, pages 3304–3311. IEEE, 2011. 1

[2] Sulabh Kumra, Shirin Joshi, and Ferat Sahin. Antipodal robotic grasping using generative residual convolutional neural network. In *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 9626–9633. IEEE, 2020. 1

[3] Douglas Morrison, Peter Corke, and Jürgen Leitner. Closing the loop for robotic grasping: A real-time, generative grasp synthesis approach. *arXiv preprint arXiv:1804.05172*, 2018. 1

[4] Ethan Perez, Florian Strub, Harm De Vries, Vincent Dumoulin, and Aaron Courville. FiLM: Visual reasoning with a general conditioning layer. In *Proceedings of the AAAI conference on artificial intelligence*, 2018. 1

[5] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pages 8748–8763. PmLR, 2021. 1

[6] Mohit Shridhar, Lucas Manuelli, and Dieter Fox. CLIPort: What and where pathways for robotic manipulation. In *Conference on robot learning*, pages 894–906. pmlr, 2022. 1

[7] Georgios Tziafas, Yucheng Xu, Arushi Goel, Mohammadreza Kasaei, Zhibin Li, and Hamidreza Kasaei. Language-guided robot grasping: CLIP-based referring grasp synthesis in clutter. *arXiv preprint arXiv:2311.05779*, 2023. 1

[8] An Dinh Vuong, Minh Nhat Vu, Baoru Huang, Nghia Nguyen, Hieu Le, Thieu Vo, and Anh Nguyen. Language-driven grasp detection. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17902–17912. IEEE, 2024. 1, 2

[9] An Dinh Vuong, Minh Nhat Vu, Hieu Le, Baoru Huang, Huynh Thi Thanh Binh, Thieu Vo, Andreas Kugi, and Anh Nguyen. Grasp-anything: Large-scale grasp dataset from foundation models. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pages 14030–14037. IEEE, 2024. 1

[10] Chong Zhou, Chen Change Loy, and Bo Dai. Extract free dense labels from CLIP. In *European conference on computer vision*, pages 696–712. Springer Nature Switzerland, 2022. 1